

Intro to Neural Networks & Deep Learning (ECS659U/ECS659P)

Yorgos Tzimiropoulos

About me :)

- Studied Electrical & Computer Engineering in Greece
- MSc & PhD at Imperial College London
- Currently:
 - Principal Scientist @ Samsung AI Center, Cambridge
 - Associate Professor @ QMUL
- Area of expertise: Computer Vision & Deep Learning



Intro to Neural Networks & Deep Learning

- What AI is all about?
- What kind of AI this module will be focusing on
- A simple example of the models we will be working on
- Module logistics

What is meant by Human Intelligence?

- Exams, education, perform well in intelligence (IQ) tests
- Skills – languages, chess, mathematical reasoning
- Cognition: perception, memory, judgement and reasoning
- Volition: the will, purpose, choice
- Emotion: affective intelligence
- Speed, efficiency and creativity of information processing

How far are we from building machines with
this type of intelligence?

General vs. Narrow Artificial Intelligence

- Gary Marcus: “*Today’s AI has been very successful in building systems that perform a single narrow goal extremely well*”
 - E.g. chess playing or identifying dog breeds
 - Solving pattern recognition
- Still this is both extremely useful and exciting!!
 - Major industry players (Google, Microsoft, IBM, Apple, Facebook, Samsung etc.) are investing Billions on AI
- The focus of this module

Applications: Computer Vision



Image Understanding



Image Generation

Applications: Face Analysis

3D Face Reconstruction from a Single Image



Download Wavefront OBJ File
(colours are stored per-vertex)

Try another image

Z Translate:

Show background image

 Tweet  Star 4,273

<https://cvl-demos.cs.nott.ac.uk/vrn/>

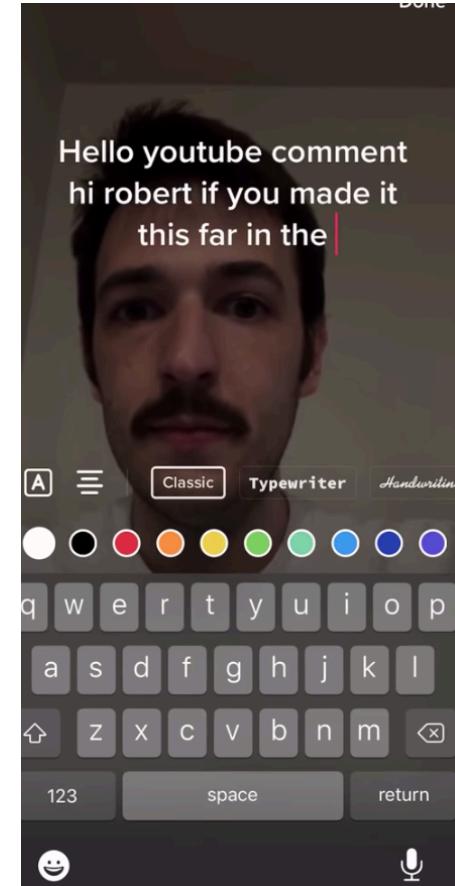
Applications: Speech Recognition



Alexa



Siri



Text to Speech
(TikTok)

Applications: Natural Language Processing

- I: Jane went to the hallway.
- I: Mary walked to the bathroom.
- I: Sandra went to the garden.
- I: Daniel went back to the garden.
- I: Sandra took the milk there.
- Q: Where is the milk?
- A: garden

Chat bot: Question Answering



Another attempt at a longer piece. An imaginary Jerome K. Jerome writes about Twitter. All I seeded was the title, the author's name and the first "It", the rest is done by #gpt3

Here is the full-length version as a PDF:
drive.google.com/file/d/1qtPa1c...

The importance of being on twitter

by Jerome K. Jerome
London, Summer 1897

It is a curious fact that the last remaining form of social life in which the people of London are still interested is Twitter. I was struck with this curious fact when I went on one of my periodical holidays to the sea-side, and found the whole place twittering like a starling-cage. I called it an anomaly, and it is.

I spoke to the sexton, whose cottage, like all sexton's cottages, is full of antiquities and interesting relics of former centuries. I said to him, "My dear sexton, what does all this twittering mean?" And he replied, "Why, sir, of course it means Twitter." "Ah!" I said, "I know about that. But what is Twitter?"

GPT-3

Current state of AI

All of these amazing technologies (and many more!) are produced by a single technology:

Deep Learning

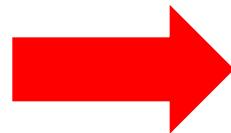
Deep Learning: The model

- Very conveniently, Deep Learning models (general speaking) follow a single pattern/paradigm
- This is the model we will be heavily using throughout the module
- This is great because if we understand the workings of this model we can apply it to many other problems, too

Deep Learning: The model

Input Data

- Images
- Speech/Audio
- Words
- Tabular data



Prediction

- A car or a plane?
- French or English?
- Answer to a question
- Price of a house

The Model

- Has some model parameters
- Uses the parameters to compute a function/combination of the input data in order to make predictions

Example

Input Data



Model



Prediction

[height, weight, age, exerc_hours]

blood_pressure

```
# [height, weight, age, exerc_hours]
# Assume 3 different subjects
X = [[180, 89, 35, 1],
      [160, 49, 40, 4],
      [170, 69, 20, 2]]
```

```
# Model W (a vector)
# W = [w1, w2, w3, w4]
W = [0.3, 0.8, -0.4, 10]
```

```
# Model output
def model_out(W, x):
    out = W[0]*x[0] + W[1]*x[1]
    + W[2]*x[2] + W[3]*x[3]
    return out
```

```
# prediction
x = X[0]
print(model_out(W, x))
```

125.2

Training the Model

- Training the model means finding the values of the model parameters
- We need training data: input data for which we know the real outputs

[height, weight, age, exerc_hours]

```
[[10, 30, 40, 1],  
 [80, 70, 60, 2],  
 [40, 10, 80, 6],  
 [50, 0, 70, 0],  
 [50, 90, 0, 1],  
 [60, 50, 60, 3],  
 [20, 50, 40, 1],  
 [40, 80, 70, 2],  
 [40, 70, 70, 5],  
 [50, 0, 60, 2],  
 [80, 30, 80, 3],  
 [90, 40, 50, 1],  
 [60, 90, 90, 6],  
 [80, 90, 90, 6],  
 [80, 50, 60, 1],  
 [0, 70, 10, 6],  
 [70, 70, 90, 8],  
 [50, 90, 0, 4],  
 [10, 80, 50, 1],  
 [10, 40, 70, 2]]
```

blood_pressure

```
[[160],  
 [40],  
 [120],  
 [180],  
 [100],  
 [0],  
 [20],  
 [50],  
 [40],  
 [80],  
 [110],  
 [110],  
 [120],  
 [130],  
 [150],  
 [170],  
 [140],  
 [100],  
 [130],  
 [140]]
```

Training the Model

To train the model

- Start from some random values for the model parameters
- Update the parameters so that for each input the model predictions are “as close as possible” to the corresponding real output values
- “as close as possible:”

```
# y is the real output,  
# and y_hat the model prediction  
def loss(y, y_hat):  
    return (y - y_hat)**2
```

$$\text{loss} = (y - \hat{y})^2$$

$$\text{loss} = (y - W[0] * x[0] + W[1] * x[1] + W[2] * x[2] + W[3] * x[3])^2$$

Training the Model

To train the model

- After we calculate the loss, each model parameter is updated using the derivative of the loss with respect to that parameter.
- E.g. for the first model parameter

$$w[0] \leftarrow w[0] - \frac{\theta loss}{\theta W[0]}$$

- The only difficult part in this algorithm could be the calculation of the derivative which requires mathematical calculations
- Fortunately, PyTorch calculates this for us automatically!

Syllabus

Week 01: Intro to AI and Tensor Data Processing

Week 02: Automatic Differentiation

Week 03: Linear Regression with Automatic Differentiation

Week 04: Softmax Regression: single-layer Neural Networks

Week 05: Multilinear Perceptrons (MLPs)

Week 06: Introduction to Assignment & Catching-up

Week 07: Convolutional Neural Networks (CNNs)

Week 08: Advanced CNNs and Optimization Algorithms

Week 09: Introduction to Natural Language Processing (NLP)

Week 10: Recurrent Neural Networks (RNNs)

Week 11: Advanced RNNs: GRUs and LSTMs

Week 12: Revision

How do I study for this module?

- Slides cover deep learning theory and PyTorch code
- Study the slides to understand the concepts
- Run the provided Jupyter Notebooks on your own!
 - Go inside the code and check the variables (print the variables, their dimensions etc.), make sure you understand what the functions do
- We cannot cover all PyTorch documentation in the lecture. When unsure about something check:
 - Official webpage: <https://pytorch.org>
 - Tutorial section: <https://pytorch.org/tutorials/>
 - PyTorch forums: <https://discuss.pytorch.org>
 - Google!

Module Logistics

- Lecture every Friday 09:00-11:00
- Lab every Friday 11:00-13:00
- Coursework
 - One item
 - Done individually
 - 50% of module total
- Examination in May:
 - 50% of module total

Recommended Books

- Aston Zhang, Zachary C. Lipton, Mu Li and Alexander J. Smola, **Dive into Deep Learning**, 2020
 - The textbook used to develop the materials
 - <https://d2l.ai>
 - This book is really amazing!
- Ian Goodfellow, Yoshua Bengio and Aaron Courville, **Deep Learning**, MIT Press, 2016
 - www.deeplearningbook.org