

**ECS7024 Statistics for Artificial Intelligence and Data
Science**

**Topic 5: Discrete
Distributions**

William Marsh

Outline

- Aims:
 - Introduce discrete probability distributions
- Coin flipping: Bernoulli distribution
- Multiple flips: the i. i. d. assumption
- Binomial distribution

Flip a Coin

Flip a Coin

- Probability of heads is 50%
 - Provided coin is fair
- Generalise: $\text{flip}(p)$
 - p is probability

$$y_1, y_2, y_3, y_{41} \sim \text{bernoulli}(0.5)$$

```
p = 0.5  
rv = stats.bernoulli(p)  
ys = rv.rvs(4)
```



Aside: random numbers from computers

Repeatedly Flipping

- Repeatedly (10 times) try 4 flips

```
p = 0.5
rv = stats.bernoulli(p)
for x in range(0, 10):
    print(rv.rvs(4))
```

```
[0 0 1 0]
[1 0 1 1]
[0 1 0 1]
[0 0 1 0]
[0 1 1 0]
[0 1 0 0]
[0 0 0 1]
[1 1 1 1]
[1 1 0 0]
[0 1 1 1]
```

Repeatedly Flipping (p=0.3)

- Repeatedly (10 times) try 4 flips

```
p = 0.3
rv = stats.bernoulli(p)
for x in range(0, 10):
    print(rv.rvs(4))
```

```
[0 0 0 0]
[0 0 0 0]
[0 1 0 1]
[0 0 0 1]
[1 1 1 0]
[1 1 1 1]
[0 1 0 0]
[1 1 1 0]
[0 0 0 0]
[0 1 0 1]
```

Probability of Outcome

- Suppose p is 0.3
 - Result of 4 flips: 1, 0, 0, 1
 - Probability?

$$\begin{aligned}Pr(1,0,0,1) &= 0.3 \times 0.7 \times 0.7 \times 0.3 \\ &= 0.3^2 \times 0.7^2\end{aligned}$$

- Probability does not depend on order
 - Assumed independence (i.i.d)
 - Only count of '1' outcomes

Probably of Outcome II

- Any p , any number of flips N

$$Pr(n \text{ from } N \mid p) = p^n \cdot (1 - p)^{N-n}$$

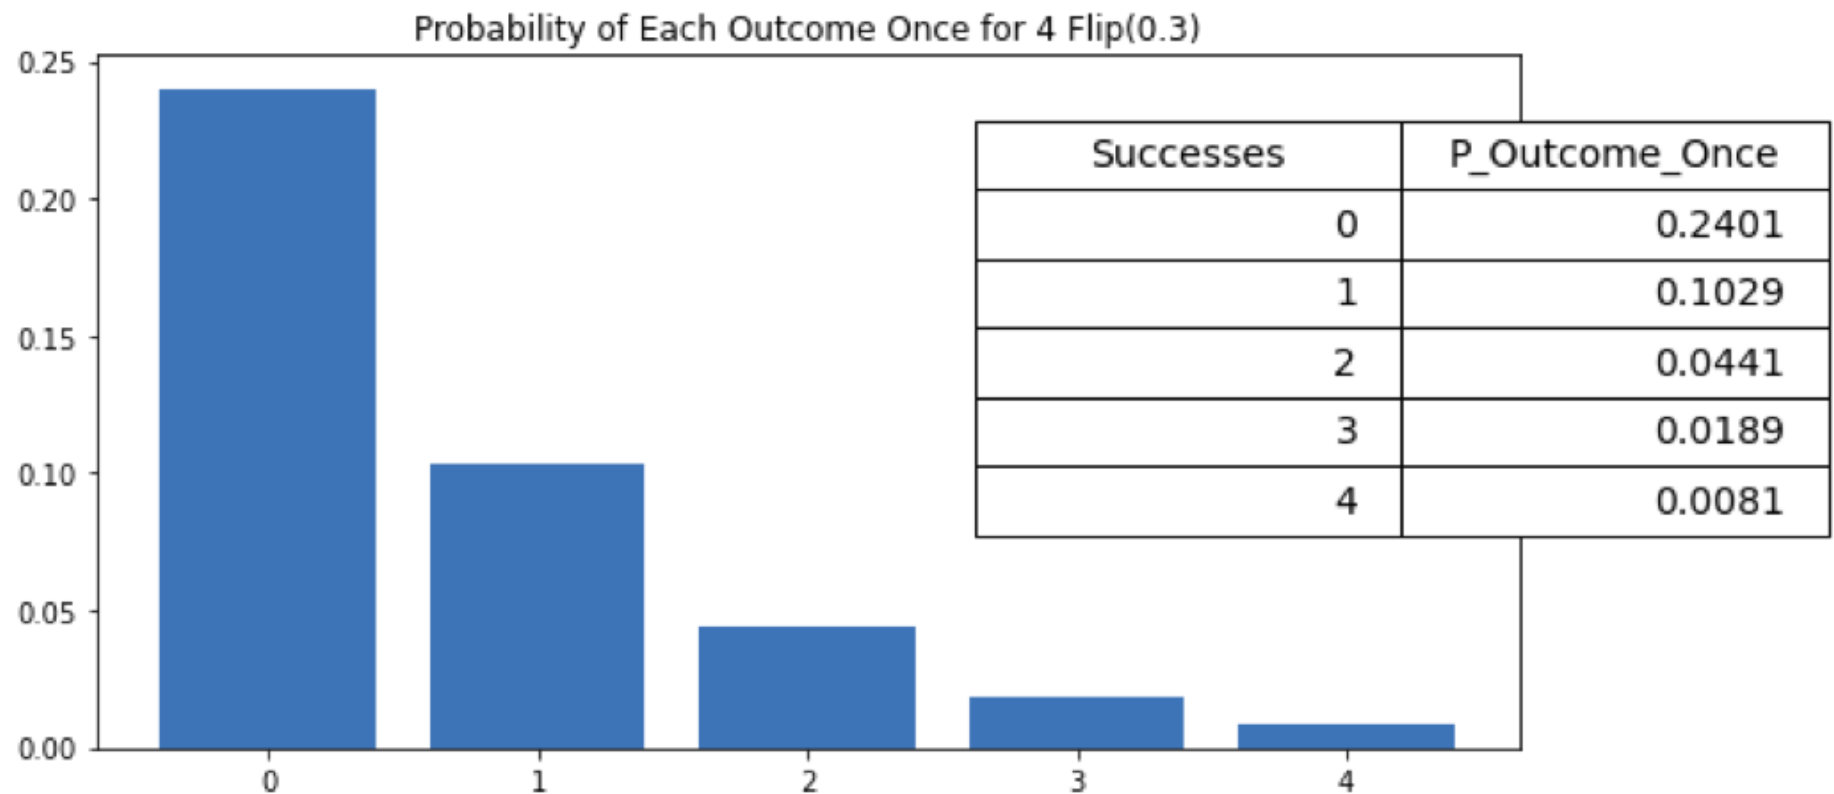
- Sufficient statistic
 - Number of positive 'n' outcomes from 'N' trials

Binomial Distribution

Probability distribution for outcomes of
coin flipping (Bernoulli) trials

Probabilities of Outcomes of 4 Flips


- Calculate probabilities of all outcomes



- But not a probability distribution

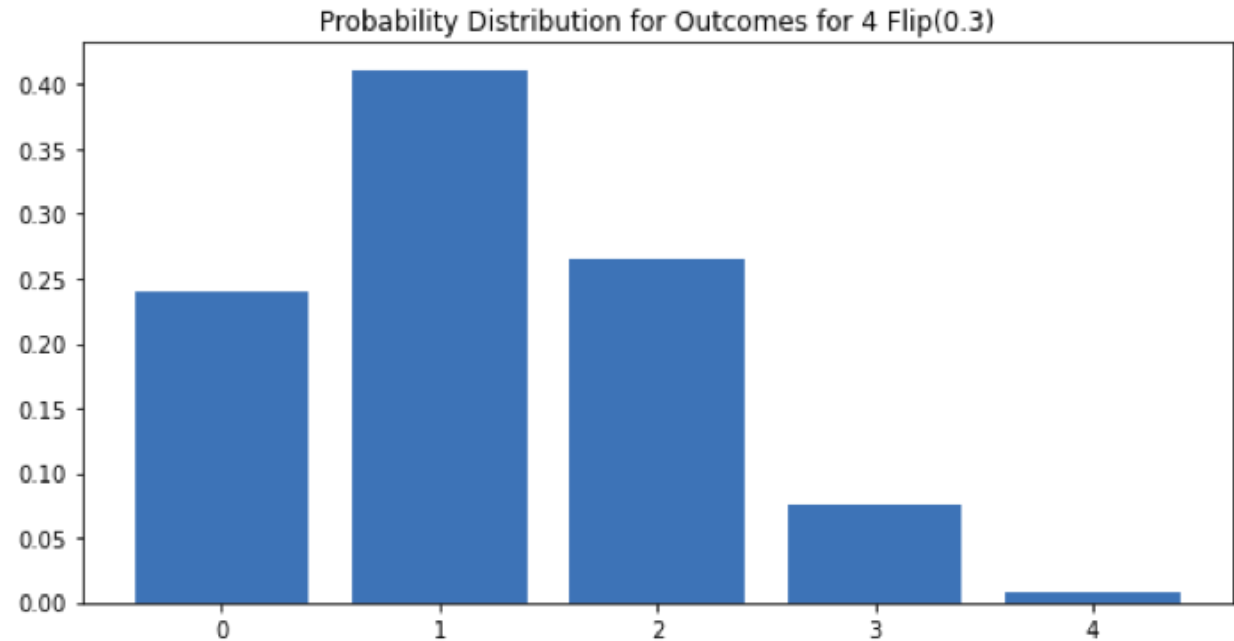
A Trial of 4 Flips

- 16 possible outcomes (incl order)
- 5 different 'number of successes'
 - Two extremes outcomes (0 or 4 successes) occur once
 - Two success occur more often
- Intuition: 2 heads and 2 tails of a coin more likely



Probability Distribution of Outcomes

- Now a probability distribution
- 'Binomial'



n successes

Pr of one outcome with n successes


Number of outcomes with n successes

Successes	P_Outcome_Once	Combinations	P_Outcome
0	0.2401	1.0	0.2401
1	0.1029	4.0	0.4116
2	0.0441	6.0	0.2646
3	0.0189	4.0	0.0756
4	0.0081	1.0	0.0081

Outcome probability

Binomial Distribution

- Distribution of outcome of a binary variable
 - Number of 'trials' (e.g. coin flips) – N
 - Probability of success (e.g. 50%) – p



`binom(4, 0.5)`

Sampling

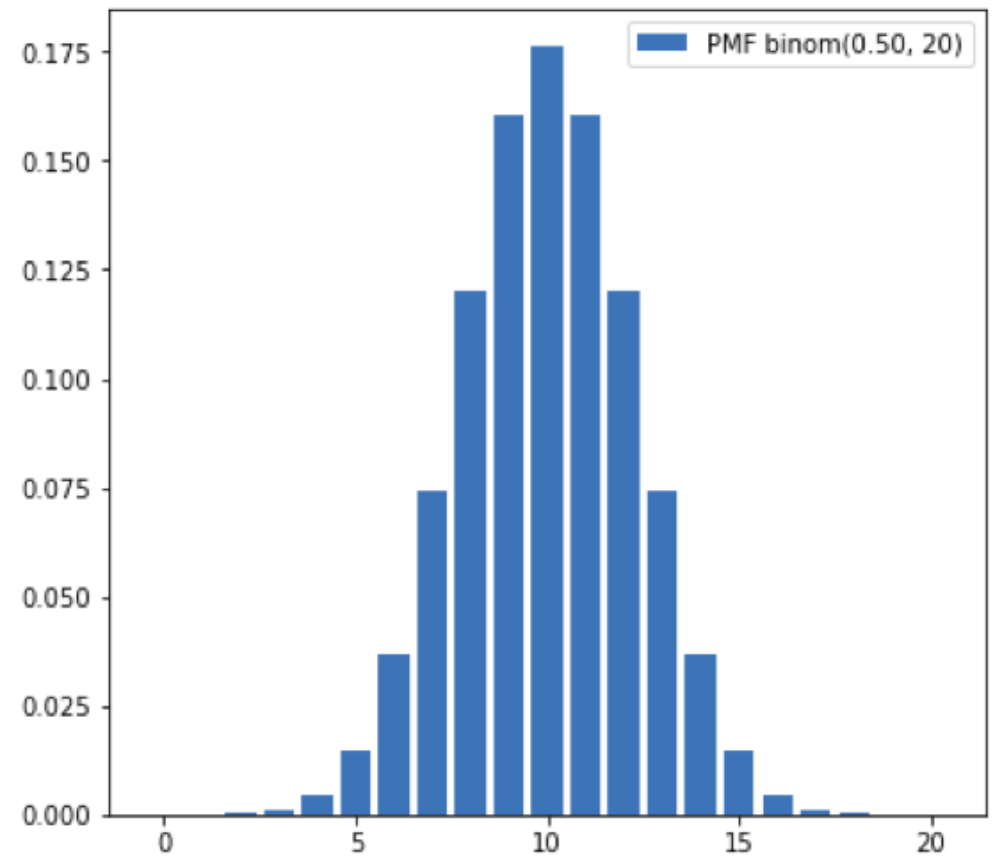

- Flip a coin – sample from Bernoulli(0.5)

$$\textit{coin} \sim \textit{Bernoulli}(0.5)$$

- Run a Bernoulli trial, 5 flips – sample from Binomial

$$\textit{successes} \sim \textit{Binomial}(5, 0.5)$$

50 Samples from Binom(20, 0.5)



Quiz 1

Every lecture will have a 'learning reflection' slide

Understanding Selection

```
.loc[]
```

Selection

- The general form of the `.loc` call is:

- `.loc[row selector, column selector]`

Use : for
all rows

The column selector is
ignored if it is omitted

- The 'selectors' can be either:
 - a value or list of values
 - an expression that is true or false

Index!!

expression such as
`df.Area == 'Tower Hamlets'`
Evaluates to a series of True / False

Binomial Probability Mass Function

Binomial PMF

- For a trial of 4 flips, there are 5 possible outcomes
 - 0, 1, 2, 3, 4 successes
 - Distribution shows the probability of each outcome
- Probability of n successes in a trial of N flips, with success probability p :

Num of trial
outcomes
with n
successes

$$\binom{N}{n} p^n (1 - p)^{(N-n)}$$

Prob of n successes

Prob of $N-n$ failures

`N= 10`

`p = 0.5`

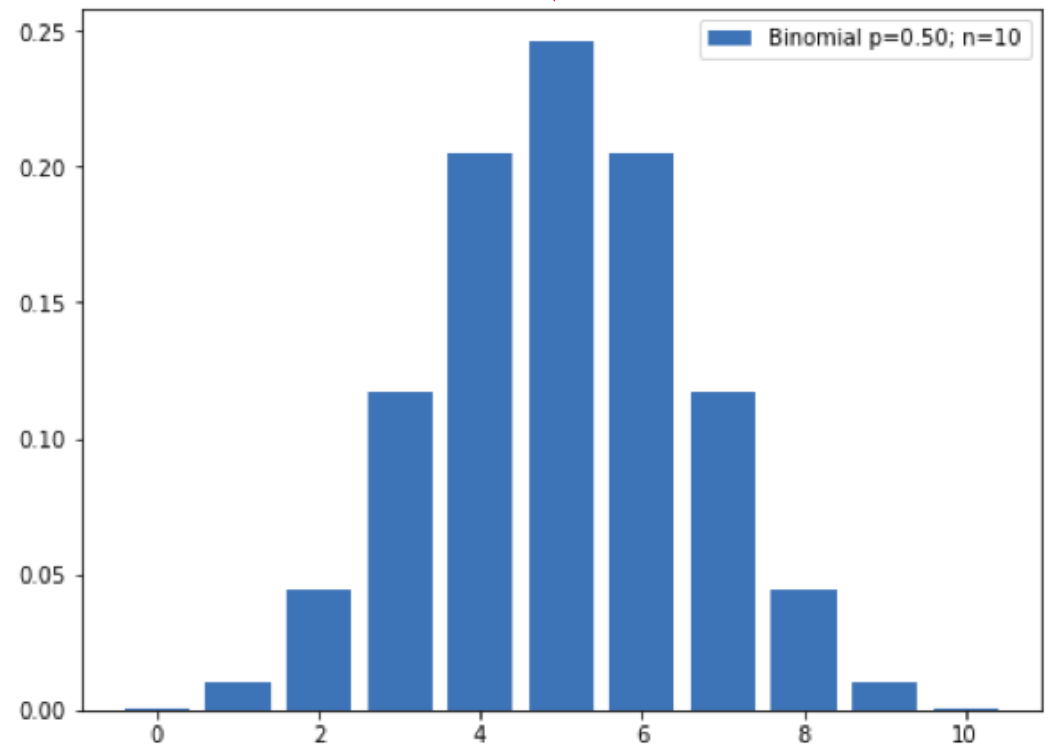
`stats.binom(N, p).pmf(3)`

Expected Value and Skew

Expected Value

- Simple formula for expected (average) number of successes
 - Trial of N flips
 - Probability of success p
 - $E[n] = N.p$


$N = 10$
 $p = 0.5$
 $E[n] = 5$



Skew of Binomial Distribution

- Skew depends on n and p
- Case 1:
expected
value
close to
 $N/2$


$N = 10$
 $p = 0.5$
 $E[n] = 5$



Skew of Binomial Distribution

- Skew depends on n and p
- Case 2:
expected
value
close to 0


$N = 10$
 $p = 0.05$
 $E[n] = 0.5$



Skew of Binomial Distribution

- Skew depends on n and p
- Case 3:
expected
value
close to N

$N = 10$
 $p = 0.95$
 $E[n] = 9.5$



Poisson Distribution

What if N is not limited?

- A hospital serves a large population
- Most people stay well
- We collect data on the average rate (people / day) of admission to hospital
- What is the distribution of daily admissions?

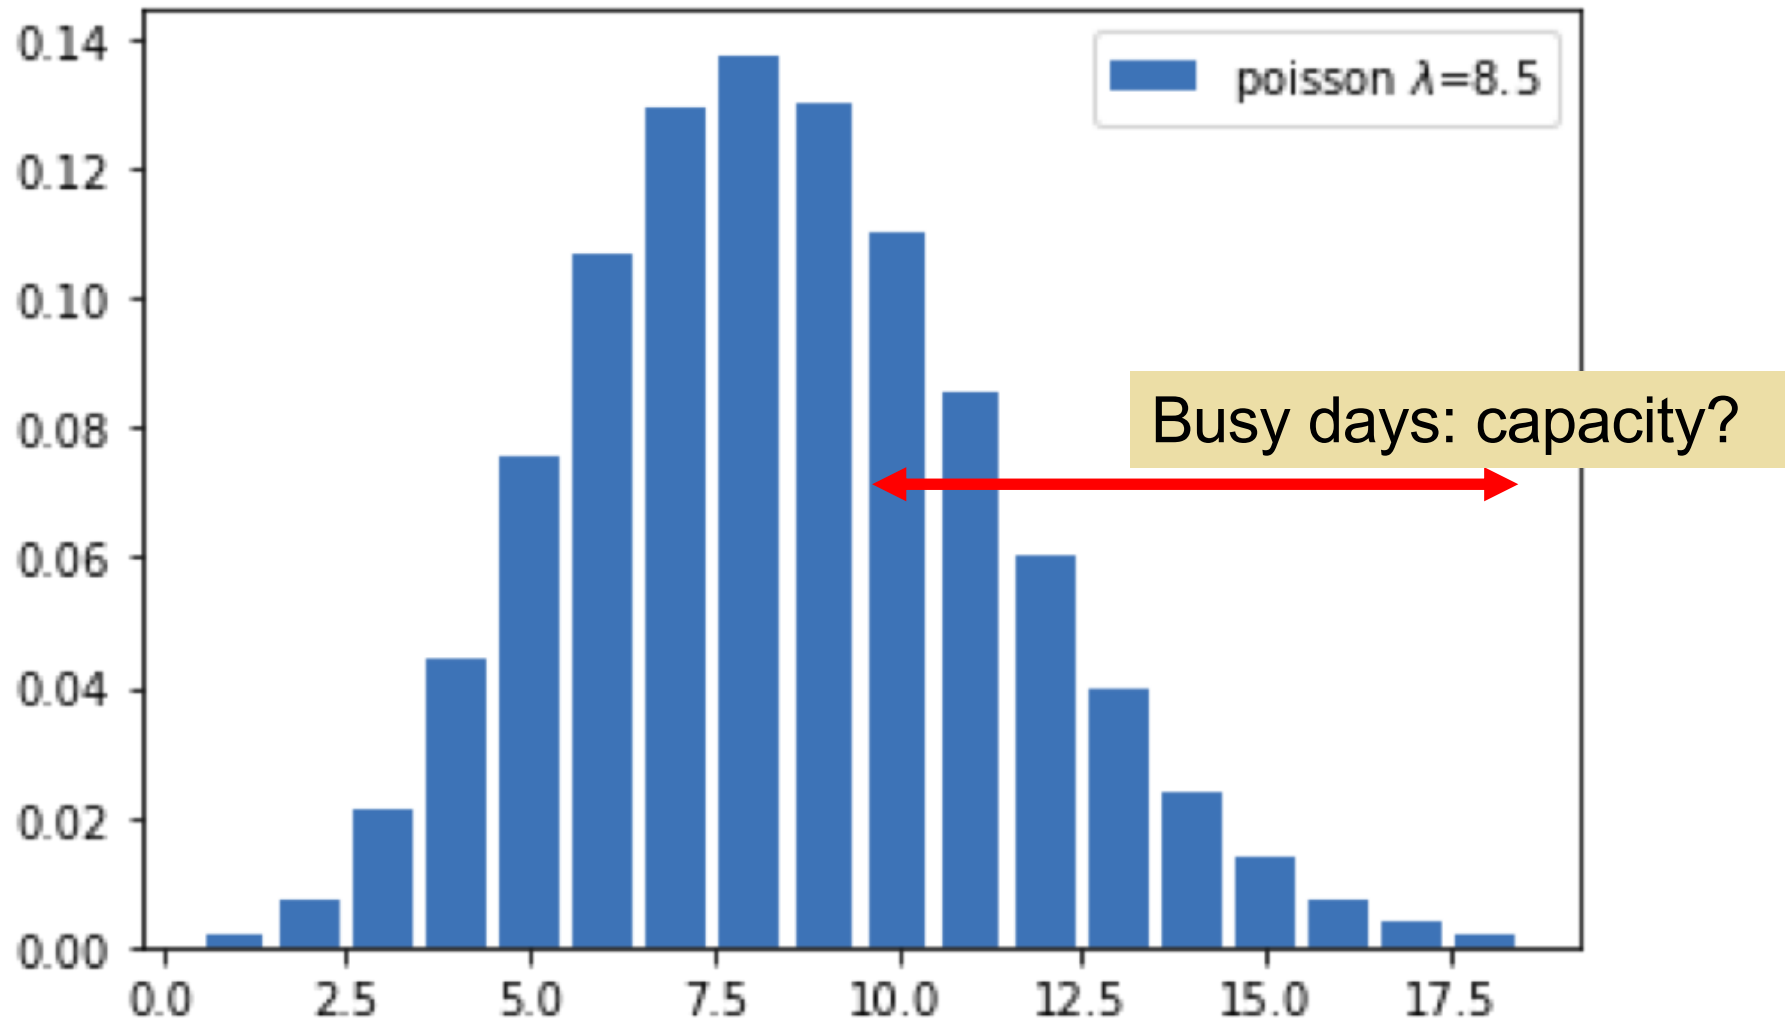
$$x \sim \text{Poisson}(\lambda)$$

Integer result:
0 upwards

Rate: real
number \geq zero

Poisson Example

- Expected value = λ



Quiz 2

Summary

Probability Distribution

- A mathematical formula describing a probability distribution
- Each corresponds to different assumptions about how variation arises
- Some common distributions

Name	Type	Description
categorical(p s)	Discrete	Generalisation of the Bernoulli to >2 outcomes
binomial(n, p)	Discrete	Outcome of n trials, which probability p
poisson(λ)	Discrete	Number of events, occurring at average rate λ
normal(μ, σ)	Continuous	Measurements with mean μ and standard deviation σ
exponential(λ)	Continuous	Time between events, occurring at average rate λ

Summary

- Bernoulli (or flip) distribution
 - Binary outcome – probability p
 - Simplest of all
- Binomial distribution
 - First example of a parametric distribution
 - Result of a Bernoulli trial (flipping repeatedly)
- Idea of sampling